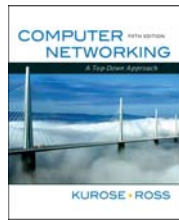


## Chapter 2 Application Layer



*Computer Networking:  
A Top Down Approach,  
5th edition.*  
Jim Kurose, Keith Ross  
Addison-Wesley, April  
2009.

All material copyright 1996-2009  
J.F. Kurose and K.W. Ross, All Rights Reserved

2: Application Layer 1

## Chapter 2: Application Layer

### Our goals:

- ❑ conceptual, implementation aspects of network application protocols
  - ❖ transport-layer service models
  - ❖ client-server paradigm
  - ❖ peer-to-peer paradigm
- ❑ learn about protocols by examining popular application-level protocols
  - ❖ HTTP
  - ❖ FTP
  - ❖ SMTP / POP3 / IMAP
  - ❖ DNS
- ❑ programming network applications
  - ❖ socket API

2: Application Layer 2

## Some network apps

- ❑ e-mail
- ❑ web
- ❑ instant messaging
- ❑ remote login
- ❑ P2P file sharing
- ❑ multi-user network games
- ❑ streaming stored video clips
- ❑ voice over IP
- ❑ real-time video conferencing
- ❑ grid computing

2: Application Layer 3

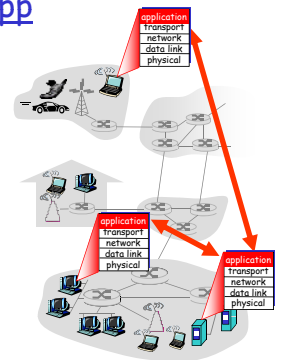
## Creating a network app

### write programs that

- ❖ run on (different) end systems
- ❖ communicate over network
- ❖ e.g., web server software communicates with browser software

### No need to write software for network-core devices

- ❖ Network-core devices do not run user applications
- ❖ applications on end systems allows for rapid app development, propagation



2: Application Layer 4

## Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 FTP
- ❑ 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- ❑ 2.5 DNS
- ❑ 2.6 P2P applications
- ❑ 2.7 Socket programming with TCP
- ❑ 2.8 Socket programming with UDP
- ❑ 2.9 Building a Web server

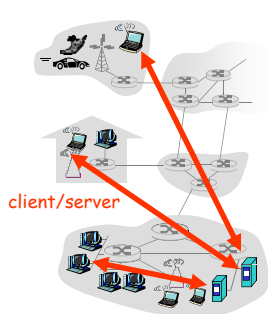
2: Application Layer 5

## Application architectures

- ❑ Client-server
- ❑ Peer-to-peer (P2P)
- ❑ Hybrid of client-server and P2P

2: Application Layer 6

## Client-server architecture



### server:

- ❖ always-on host
- ❖ permanent IP address
- ❖ server farms for scaling

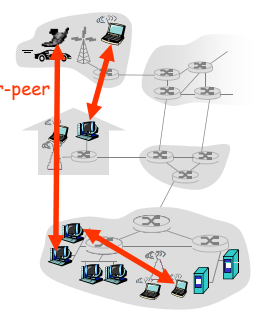
### clients:

- ❖ communicate with server
- ❖ may be intermittently connected
- ❖ may have dynamic IP addresses
- ❖ do not communicate directly with each other

2: Application Layer 7

## Pure P2P architecture

- ❑ no always-on server
- ❑ arbitrary end systems directly communicate
- ❑ peers are intermittently connected and change IP addresses



Highly scalable but difficult to manage

2: Application Layer 8

## Hybrid of client-server and P2P

### Skype

- ❖ voice-over-IP P2P application
- ❖ centralized server: finding address of remote party:
- ❖ client-client connection: direct (not through server)

### Instant messaging

- ❖ chatting between two users is P2P
- ❖ centralized service: client presence detection/location
  - user registers its IP address with central server when it comes online
  - user contacts central server to find IP addresses of buddies

2: Application Layer 9

## Processes communicating

**Process:** program running within a host.

- ❑ within same host, two processes communicate using **inter-process communication** (defined by OS).
- ❑ processes in different hosts communicate by exchanging **messages**

**Client process:** process that initiates communication

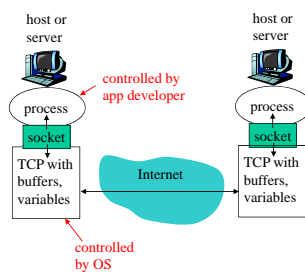
**Server process:** process that waits to be contacted

- ❑ Note: applications with P2P architectures have client processes & server processes

2: Application Layer 10

## Sockets

- ❑ process sends/receives messages to/from its **socket**
- ❑ socket analogous to door
  - ❖ sending process shoves message out door
  - ❖ sending process relies on transport infrastructure on other side of door which brings message to socket at receiving process
- ❑ API: (1) choice of transport protocol; (2) ability to fix a few parameters (lots more on this later)



2: Application Layer 11

## Addressing processes

- ❑ to receive messages, process must have **identifier**
- ❑ host device has unique 32-bit IP address
- ❑ **Q:** does IP address of host suffice for identifying the process?

2: Application Layer 12

## Addressing processes

- to receive messages, process must have **identifier**
- host device has unique 32-bit IP address
- Q:** does IP address of host on which process runs suffice for identifying the process?
  - A:** No, many processes can be running on same host
- identifier** includes both **IP address** and **port numbers** associated with process on host.
- Example port numbers:
  - HTTP server: 80
  - Mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - IP address:** 128.119.245.12
  - Port number:** 80
- more shortly...

2: Application Layer 13

## App-layer protocol defines

- Types of messages exchanged,
    - e.g., request, response
  - Message syntax:
    - what fields in messages & how fields are delineated
  - Message semantics
    - meaning of information in fields
  - Rules for when and how processes send & respond to messages
- Public-domain protocols:**
- defined in RFCs
  - allows for interoperability
  - e.g., HTTP, SMTP
- Proprietary protocols:**
- e.g., Skype

2: Application Layer 14

## What transport service does an app need?

### Data loss

- some apps (e.g., audio) can tolerate some loss
- other apps (e.g., file transfer, telnet) require 100% reliable data transfer

### Timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be "effective"

### Throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be "effective"
- other apps ("elastic apps") make use of whatever throughput they get

### Security

- Encryption, data integrity, ...

2: Application Layer 15

## Transport service requirements of common apps

Application	Data loss	Throughput	Time Sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100's msec
interactive games	loss-tolerant	few kbps up	yes, 100's msec
instant messaging	no loss	elastic	yes and no

2: Application Layer 16

## Internet transport protocols services

### TCP service:

- connection-oriented:** setup required between client and server processes
- reliable transport:** between sending and receiving process
- flow control:** sender won't overwhelm receiver
- congestion control:** throttle sender when network overloaded
- does not provide:** timing, minimum throughput guarantees, security

### UDP service:

- unreliable data transfer between sending and receiving process
- does not provide: connection setup, reliability, flow control, congestion control, timing, throughput guarantee, or security

**Q:** why bother? Why is there a UDP?

2: Application Layer 17

## Internet apps: application, transport protocols

Application	Application layer protocol	Underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (eg Youtube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	typically UDP

2: Application Layer 18

## Chapter 2: Application layer

- 2.1 Principles of network applications
  - ❖ app architectures
  - ❖ app requirements
- 2.2 Web and HTTP
- 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP

## Web and HTTP

### First some jargon

- Web page consists of **objects**
- Object can be HTML file, JPEG image, Java applet, audio file,...
- Web page consists of **base HTML-file** which includes several referenced objects
- Each object is addressable by a **URL**
- Example URL:

www.someschool.edu / someDept/pic.gif  
host name path name

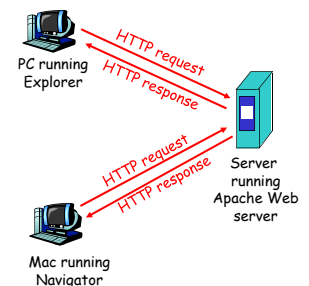
## HTML Example

```
<HTML>
<HEAD>
<TITLE> PB's HomePage </TITLE>
</HEAD>
<BODY>
<CENTER><IMG SRC = "bad_picture.gif" ALT = " " ><BR></CENTER>
<P><CENTER><H1>USC Computer Science and Engineering
Department</H1></CENTER></p>
Welcome to my goofy HomePage!
...
<A HREF = http://www.cse.sc.edu/~pb/mydogs_page.html> Spot's Page </A>
</BODY>
</HTML>
```

## HTTP overview

### HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
  - ❖ **client**: browser that requests, receives, "displays" Web objects
  - ❖ **server**: Web server sends objects in response to requests



## HTTP overview (continued)

### Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

### HTTP is "stateless"

- server maintains no information about past client requests

aside  
Protocols that maintain "state" are complex!  
□ past history (state) must be maintained  
□ if server/client crashes, their views of "state" may be inconsistent, must be reconciled

## HTTP connections

### Nonpersistent HTTP

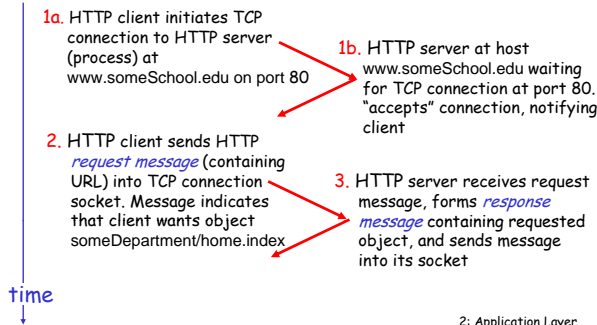
- At most one object is sent over a TCP connection.

### Persistent HTTP

- Multiple objects can be sent over single TCP connection between client and server.

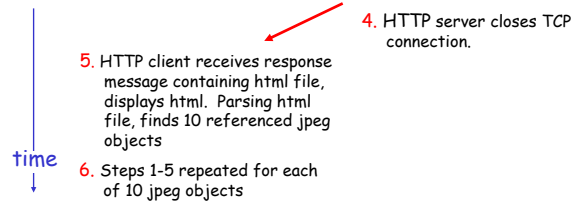
## Nonpersistent HTTP

Suppose user enters URL  
`www.someSchool.edu/someDepartment/home.index` (contains text, references to 10 jpeg images)



2: Application Layer 25

## Nonpersistent HTTP (cont.)



2: Application Layer 26

## Non-Persistent HTTP: Response time

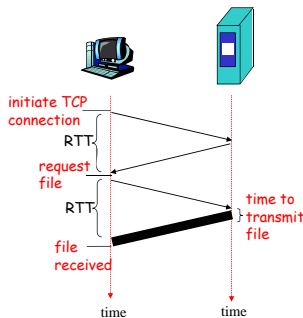
**Definition of RTT:** time for a small packet to travel from client to server and back.

**Response time:**

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return

- file transmission time

**total = 2RTT + transmit time**



2: Application Layer 27

## Persistent HTTP

**Nonpersistent HTTP issues:**

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

**Persistent HTTP**

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

2: Application Layer 28

## HTTP request message

- two types of HTTP messages: *request, response*

- HTTP request message:**

- ASCII (human-readable format)

request line (GET, POST, HEAD commands)

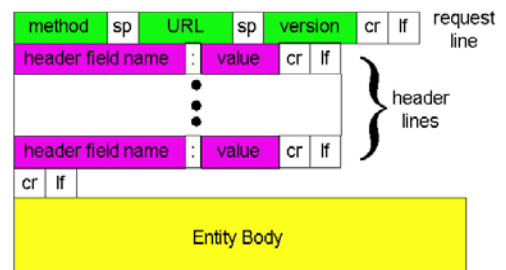
header lines

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
```

Carriage return line feed indicates end of message (extra carriage return, line feed)

2: Application Layer 29

## HTTP request message: general format



2: Application Layer 30

## Uploading form input

### Post method:

- Web page often includes form input
- Input is uploaded to server in entity body

### URL method:

- Uses GET method
- Input is uploaded in URL field of request line:

www.somesite.com/animalsearch?monkeys&banana

2: Application Layer 31

## Method types

### HTTP/1.0

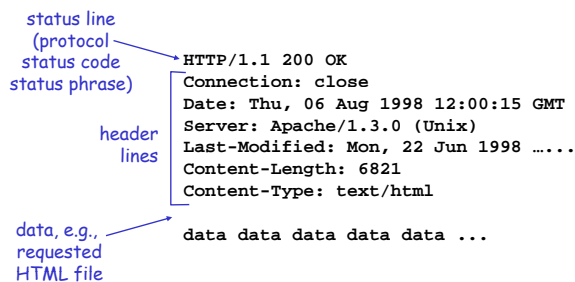
- GET
- POST
- HEAD
  - ✦ asks server to leave requested object out of response

### HTTP/1.1

- GET, POST, HEAD
- PUT
  - ✦ uploads file in entity body to path specified in URL field
- DELETE
  - ✦ deletes file specified in the URL field

2: Application Layer 32

## HTTP response message



2: Application Layer 33

## HTTP response status codes

In first line in server→client response message.

A few sample codes:

### 200 OK

- ✦ request succeeded, requested object later in this message

### 301 Moved Permanently

- ✦ requested object moved, new location specified later in this message (Location:)

### 400 Bad Request

- ✦ request message not understood by server

### 404 Not Found

- ✦ requested document not found on this server

### 505 HTTP Version Not Supported

2: Application Layer 34

## Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

```
telnet cis.poly.edu 80
```

Opens TCP connection to port 80 (default HTTP server port) at cis.poly.edu. Anything typed in sent to port 80 at cis.poly.edu

2. Type in a GET HTTP request:

```
GET /~ross/ HTTP/1.1
Host: cis.poly.edu
```

By typing this in (hit carriage return twice), you send this minimal (but complete) GET request to HTTP server

3. Look at response message sent by HTTP server!

2: Application Layer 35

## User-server state: cookies

Many major Web sites use cookies

### Four components:

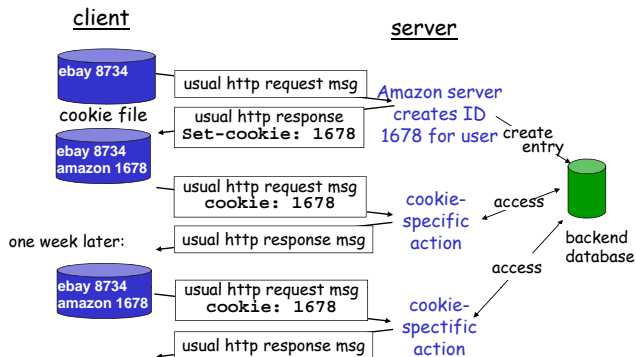
- 1) cookie header line of HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

### Example:

- Susan always access Internet always from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - ✦ unique ID
  - ✦ entry in backend database for ID

2: Application Layer 36

## Cookies: keeping "state" (cont.)



2: Application Layer 37

## Cookies (continued)

### What cookies can bring:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

### How to keep "state":

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

### aside Cookies and privacy:

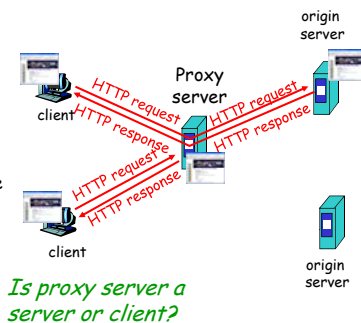
- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

2: Application Layer 38

## Web caches (proxy server)

**Goal:** satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
  - ❖ object in cache: cache returns object
  - ❖ else cache requests object from origin server, then returns object to client



2: Application Layer 39

## More about Web caching

- cache acts as both client and server
- typically cache is installed by ISP (university, company, residential ISP)

### Why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link.
- Internet dense with caches: enables "poor" content providers to effectively deliver content (but so does P2P file sharing)

2: Application Layer 40

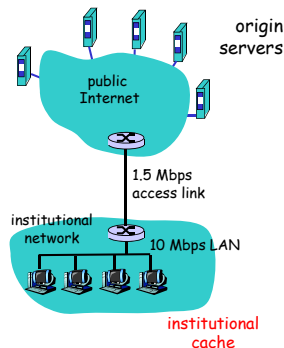
## Caching example

### Assumptions

- average object size = 100,000 bits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router = 2 sec

### Consequences

- utilization on LAN = 15%
- utilization on access link = 100%
- total delay = Internet delay + access delay + LAN delay = 2 sec + minutes + milliseconds



2: Application Layer 41

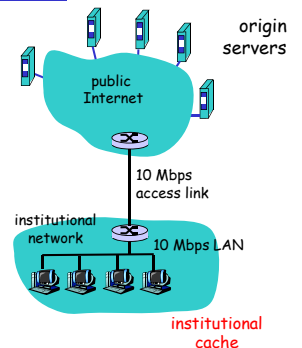
## Caching example (cont)

### possible solution

- increase bandwidth of access link to, say, 10 Mbps

### consequence

- utilization on LAN = 15%
- utilization on access link = 15%
- Total delay = Internet delay + access delay + LAN delay = 2 sec + msec + msec
- often a costly upgrade



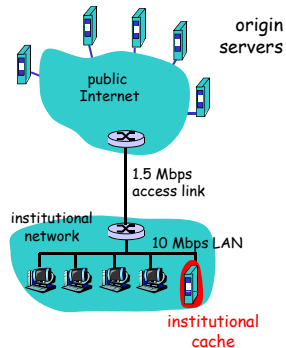
2: Application Layer 42



## Caching example (cont)

### possible solution: install cache

- suppose hit rate is 0.4
- consequence**
  - 40% requests will be satisfied almost immediately
  - 60% requests satisfied by origin server
  - utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
  - total avg delay = Internet delay + access delay + LAN delay =  $.6 * (2.01) \text{ secs} + .4 * \text{milliseconds} < 1.4 \text{ secs}$



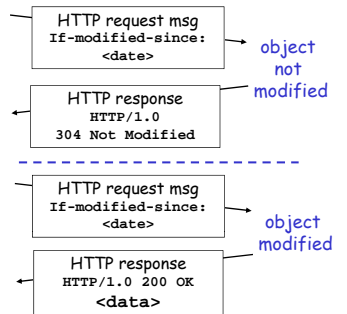
2: Application Layer 43

## Conditional GET

- Goal:** don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request  
If-modified-since: <date>
- server: response contains no object if cached copy is up-to-date:  
HTTP/1.0 304 Not Modified

cache

server



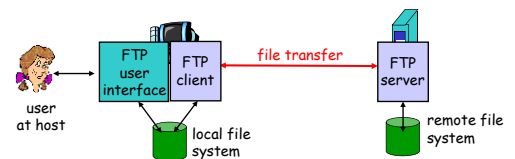
2: Application Layer 44

## Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
  - SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP
- 2.9 Building a Web server

2: Application Layer 45

## FTP: the file transfer protocol

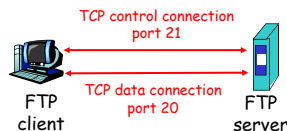


- transfer file to/from remote host
- client/server model
  - client:** side that initiates transfer (either to/from remote)
  - server:** remote host
- ftp: RFC 959
- ftp server: port 21

2: Application Layer 46

## FTP: separate control, data connections

- FTP client contacts FTP server at port 21, TCP is transport protocol
- client authorized over control connection
- client browses remote directory by sending commands over control connection.
- when server receives file transfer command, server opens 2<sup>nd</sup> TCP connection (for file) to client
- after transferring one file, server closes data connection.
- server opens another TCP data connection to transfer another file.
- control connection: "out of band"
- FTP server maintains "state": current directory, earlier authentication



2: Application Layer 47

## FTP commands, responses

### Sample commands:

- sent as ASCII text over control channel
- USER *username*
- PASS *password*
- LIST return list of file in current directory
- RETR *filename* retrieves (*gets*) file
- STOR *filename* stores (*puts*) file onto remote host

### Sample return codes

- status code and phrase (as in HTTP)
- 331 Username OK, password required
- 125 data connection already open; transfer starting
- 425 Can't open data connection
- 452 Error writing file

2: Application Layer 48



## Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
  - ✦ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP

2: Application Layer 49

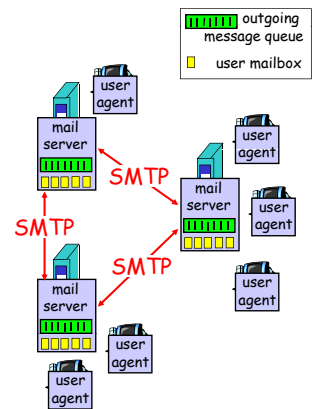
## Electronic Mail

### Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

### User Agent

- a.k.a. "mail reader"
- composing, editing, reading mail messages
- e.g., Eudora, Outlook, elm, Mozilla Thunderbird
- outgoing, incoming messages stored on server

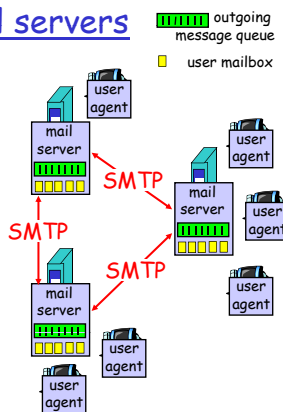


2: Application Layer 50

## Electronic Mail: mail servers

### Mail Servers

- mailbox contains incoming messages for user
- message queue of outgoing (to be sent) mail messages
- SMTP protocol between mail servers to send email messages
  - ✦ client: sending mail server
  - ✦ "server": receiving mail server



2: Application Layer 51

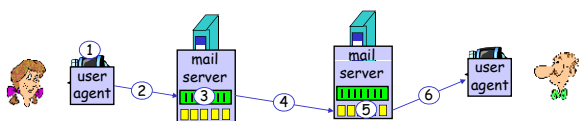
## Electronic Mail: SMTP [RFC 2821]

- uses TCP to reliably transfer email message from client to server, port 25
- direct transfer: sending server to receiving server
- three phases of transfer
  - ✦ handshaking (greeting)
  - ✦ transfer of messages
  - ✦ closure
- command/response interaction
  - ✦ commands: ASCII text
  - ✦ response: status code and phrase
- messages must be in 7-bit ASCII

2: Application Layer 52

## Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message and "to" bob@someschool.edu
- 2) Alice's UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message



2: Application Layer 53

## Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

2: Application Layer 54

## Try SMTP interaction for yourself:

- ❑ telnet servername 25
  - ❑ see 220 reply from server
  - ❑ enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands
- above lets you send email without using email client (reader)

2: Application Layer 55

## SMTP: final words

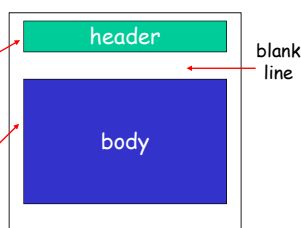
- ❑ SMTP uses persistent connections
  - ❑ SMTP requires message (header & body) to be in 7-bit ASCII
  - ❑ SMTP server uses CRLF, CRLF to determine end of message
- Comparison with HTTP:**
- ❑ HTTP: pull
  - ❑ SMTP: push
  - ❑ both have ASCII command/response interaction, status codes
  - ❑ HTTP: each object encapsulated in its own response msg
  - ❑ SMTP: multiple objects sent in multipart msg

2: Application Layer 56

## Mail message format

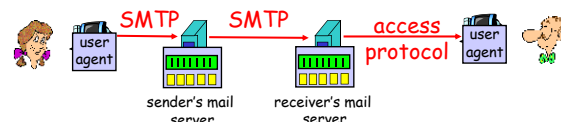
SMTP: protocol for exchanging email msgs  
RFC 822: standard for text message format:

- ❑ header lines, e.g.,
  - ✦ To:
  - ✦ From:
  - ✦ Subject:*different from SMTP commands!*
- ❑ body
  - ✦ the "message", ASCII characters only



2: Application Layer 57

## Mail access protocols



- ❑ SMTP: delivery/storage to receiver's server
- ❑ Mail access protocol: retrieval from server
  - ✦ POP: Post Office Protocol [RFC 1939]
    - authorization (agent <--> server) and download
  - ✦ IMAP: Internet Mail Access Protocol [RFC 1730]
    - more features (more complex)
    - manipulation of stored msgs on server
  - ✦ HTTP: gmail, Hotmail, Yahoo! Mail, etc.

2: Application Layer 58

## POP3 protocol

- authorization phase**
- ❑ client commands:
    - ✦ user: declare username
    - ✦ pass: password
  - ❑ server responses
    - ✦ +OK
    - ✦ -ERR

- transaction phase, client:**
- ❑ list: list message numbers
  - ❑ retr: retrieve message by number
  - ❑ dele: delete
  - ❑ quit

```

S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
  
```

2: Application Layer 59

## POP3 (more) and IMAP

### More about POP3

- ❑ Previous example uses "download and delete" mode.
- ❑ Bob cannot re-read e-mail if he changes client
- ❑ "Download-and-keep": copies of messages on different clients
- ❑ POP3 is stateless across sessions

### IMAP

- ❑ Keep all messages in one place: the server
- ❑ Allows user to organize messages in folders
- ❑ IMAP keeps user state across sessions:
  - ✦ names of folders and mappings between message IDs and folder name

2: Application Layer 60

## Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- **2.5 DNS**
- 2.6 P2P applications
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP
- 2.9 Building a Web server

2: Application Layer 61

## DNS: Domain Name System

- People:** many identifiers:
- ❖ SSN, name, passport #
- Internet hosts, routers:**
- ❖ IP address (32 bit) - used for addressing datagrams
  - ❖ "name", e.g., ww.yahoo.com - used by humans
- Q:** map between IP addresses and name ?
- Domain Name System:**
- *distributed database* implemented in hierarchy of many *name servers*
  - *application-layer protocol* host, routers, name servers to communicate to *resolve* names (address/name translation)
    - ❖ note: core Internet function, implemented as application-layer protocol
    - ❖ complexity at network's "edge"

2: Application Layer 62

## DNS

### DNS services

- hostname to IP address translation
- host aliasing
  - ❖ Canonical, alias names
- mail server aliasing
- load distribution
  - ❖ replicated Web servers: set of IP addresses for one canonical name

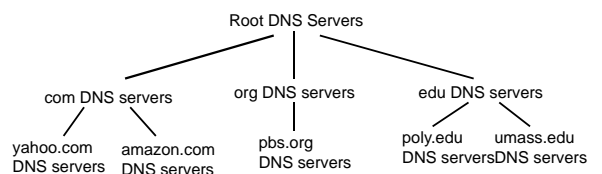
### Why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

doesn't *scale*!

2: Application Layer 63

## Distributed, Hierarchical Database



### Client wants IP for www.amazon.com; 1<sup>st</sup> approx:

- client queries a root server to find com DNS server
- client queries com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

2: Application Layer 64

## DNS: Root name servers

- contacted by local name server that can not resolve name
- root name server:
  - ❖ contacts authoritative name server if name mapping not known
  - ❖ gets mapping
  - ❖ returns mapping to local name server



2: Application Layer 65

## TLD and Authoritative Servers

- **Top-level domain (TLD) servers:**
  - ❖ responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
  - ❖ Network Solutions maintains servers for com TLD
  - ❖ Educause for edu TLD
- **Authoritative DNS servers:**
  - ❖ organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., Web, mail).
  - ❖ can be maintained by organization or service provider

2: Application Layer 66

## Local Name Server

- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one.
  - ❖ also called "default name server"
- when host makes DNS query, query is sent to its local DNS server
  - ❖ acts as proxy, forwards query into hierarchy

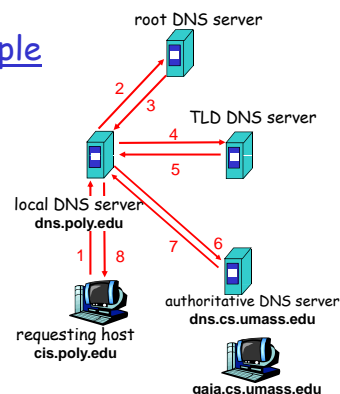
2: Application Layer 67

## DNS name resolution example

- Host at cis.poly.edu wants IP address for gaia.cs.umass.edu

### iterated query:

- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"

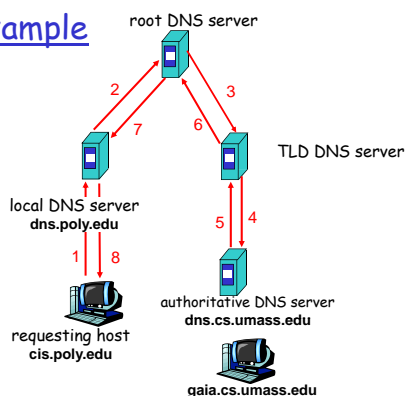


2: Application Layer 68

## DNS name resolution example

### recursive query:

- puts burden of name resolution on contacted name server
- heavy load?



2: Application Layer 69

## DNS: caching and updating records

- once (any) name server learns mapping, it *caches* mapping
  - ❖ cache entries timeout (disappear) after some time
  - ❖ TLD servers typically cached in local name servers
    - Thus root name servers not often visited
- update/notify mechanisms under design by IETF
  - ❖ RFC 2136
  - ❖ <http://www.ietf.org/html.charters/dnsind-charter.html>

2: Application Layer 70

## DNS records

**DNS:** distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

- Type=A
  - ❖ name is hostname
  - ❖ value is IP address
- Type=NS
  - ❖ name is domain (e.g. foo.com)
  - ❖ value is hostname of authoritative name server for this domain
- Type=CNAME
  - ❖ name is alias name for some "canonical" (the real) name
  - ❖ value is canonical name
  - ❖ example: www.ibm.com is really servereast.backup2.ibm.com
- Type=MX
  - ❖ value is name of mailserver associated with name

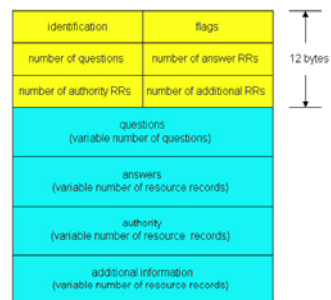
2: Application Layer 71

## DNS protocol, messages

**DNS protocol:** *query* and *reply* messages, both with same *message format*

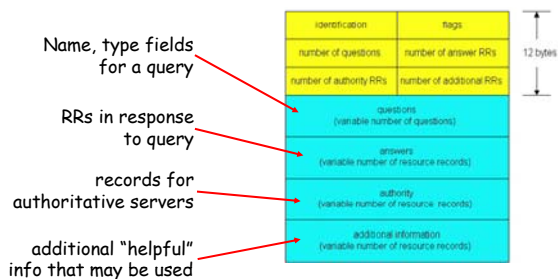
### msg header

- **identification:** 16 bit # for query, reply to query uses same #
- **flags:**
  - ❖ query or reply
  - ❖ recursion desired
  - ❖ recursion available
  - ❖ reply is authoritative



2: Application Layer 72

## DNS protocol, messages



2: Application Layer 73

## Inserting records into DNS

- example: new startup "Network Utopia"
- register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
  - ❖ provide names, IP addresses of authoritative name server (primary and secondary)
  - ❖ registrar inserts two RRs into com TLD server:
   
(networkutopia.com, dns1.networkutopia.com, NS)
   
(dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server Type A record for www.networkutopia.com; Type MX record for networkutopia.com
- How do people get IP address of your Web site?

2: Application Layer 74

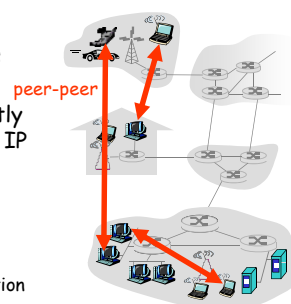
## Chapter 2: Application layer

- 2.1 Principles of network applications
  - ❖ app architectures
  - ❖ app requirements
- 2.2 Web and HTTP
- 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP

2: Application Layer 75

## Pure P2P architecture

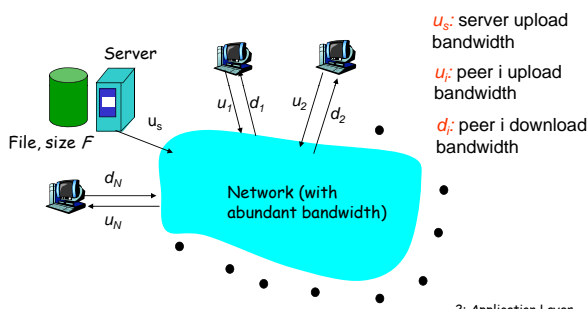
- no always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses
- Two topics:
  - ❖ File distribution
  - ❖ Searching for information



2: Application Layer 76

## File Distribution: Server-Client vs P2P

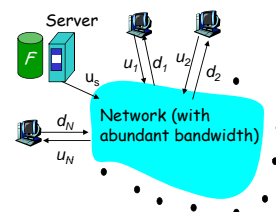
Question: How much time to distribute file from one server to  $N$  peers?



2: Application Layer 77

## File distribution time: server-client

- server sequentially sends  $N$  copies:
  - ❖  $NF/u_s$  time
- client  $i$  takes  $F/d_i$  time to download



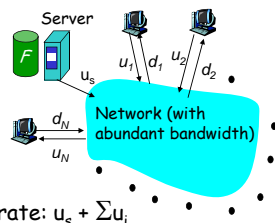
Time to distribute  $F$  to  $N$  clients using client/server approach =  $d_{cs} = \max \{ NF/u_s, F/\min_i(d_i) \}$

increases linearly in  $N$  (for large  $N$ )

2: Application Layer 78

## File distribution time: P2P

- server must send one copy:  $F/u_s$  time
- client  $i$  takes  $F/d_i$  time to download
- NF bits must be downloaded (aggregate)
  - fastest possible upload rate:  $u_s + \sum u_i$

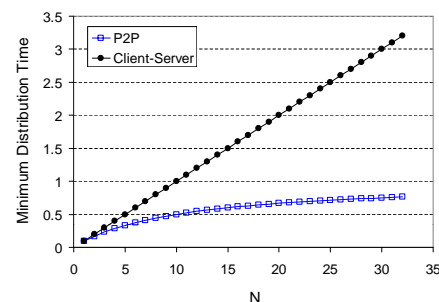


$$d_{P2P} = \max \{ F/u_s, F/\min(d_i), NF/(u_s + \sum u_i) \}$$

2: Application Layer 79

## Server-client vs. P2P: example

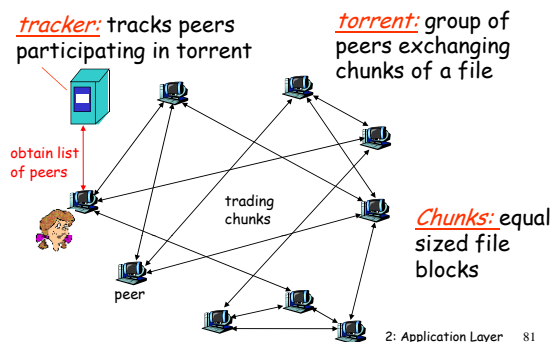
Client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{\min} \geq u_s$



2: Application Layer 80

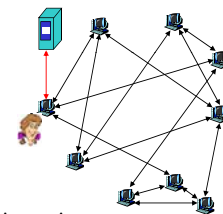
## File distribution: BitTorrent

- P2P file distribution



2: Application Layer 81

## BitTorrent (1)



- file divided into 256KB *chunks*.
- peer joining torrent:
  - has no chunks, but will accumulate them over time
  - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- while downloading, peer uploads chunks to other peers.
- peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain

2: Application Layer 82

## BitTorrent (2)

### Pulling Chunks

- at any given time, different peers have different subsets of file chunks
- periodically, a peer (Alice) asks each neighbor for list of chunks that they have.
- Alice sends requests for her missing chunks
  - rarest first

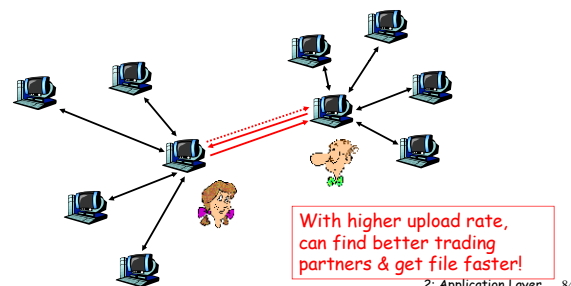
### Sending Chunks: tit-for-tat

- Alice sends chunks to four neighbors currently sending her chunks *at the highest rate*
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - newly chosen peer may join top 4
  - "optimistically unchoke"

2: Application Layer 83

## BitTorrent: Tit-for-tat

- Alice "optimistically unchokes" Bob
- Alice becomes one of Bob's top-four providers; Bob reciprocates
- Bob becomes one of Alice's top-four providers



2: Application Layer 84

## Distributed Hash Table (DHT)

- DHT = distributed P2P database
- Database has (**key**, **value**) pairs;
  - ❖ key: ss number; value: human name
  - ❖ key: content type; value: IP address
- Peers **query** DB with key
  - ❖ DB returns values that match the key
- Peers can also **insert** (key, value) peers

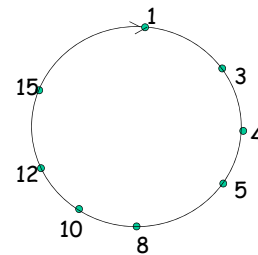
## DHT Identifiers

- Assign integer identifier to each peer in range  $[0, 2^n - 1]$ .
  - ❖ Each identifier can be represented by  $n$  bits.
- Require each key to be an integer in **same range**.
- To get integer keys, hash original key.
  - ❖ eg, key =  $h(\text{"Led Zeppelin IV"})$
  - ❖ This is why they call it a distributed "hash" table

## How to assign keys to peers?

- Central issue:
  - ❖ Assigning (key, value) pairs to peers.
- Rule: assign key to the peer that has the **closest** ID.
- Convention in lecture: closest is the **immediate successor** of the key.
- Ex:  $n=4$ ; peers: 1,3,4,5,8,10,12,14;
  - ❖ key = 13, then successor peer = 14
  - ❖ key = 15, then successor peer = 1

## Circular DHT (1)

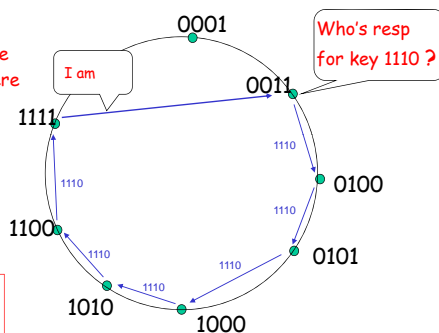


- Each peer *only* aware of immediate successor and predecessor.
- "Overlay network"

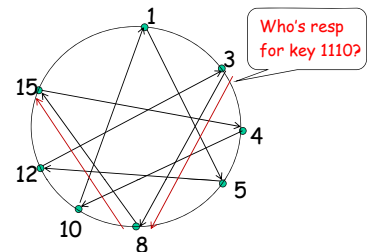
## Circle DHT (2)

$O(N)$  messages on avg to resolve query, when there are  $N$  peers

Define **closest** as closest successor



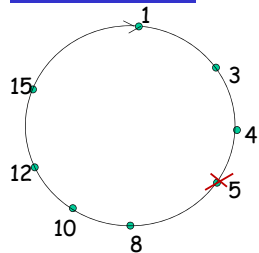
## Circular DHT with Shortcuts



- Each peer keeps track of IP addresses of predecessor, successor, short cuts.
- Reduced from 6 to 2 messages.
- Possible to design shortcuts so  $O(\log N)$  neighbors,  $O(\log N)$  messages in query



## Peer Churn



- To handle peer churn, require each peer to know the IP address of its two successors.
- Each peer periodically pings its two successors to see if they are still alive.

- Peer 5 abruptly leaves
- Peer 4 detects; makes 8 its immediate successor; asks 8 who its immediate successor is; makes 8's immediate successor its second successor.
- What if peer 13 wants to join?

## Socket programming

**Goal:** learn how to build client/server application that communicate using sockets

### Socket API

- introduced in BSD4.1 UNIX, 1981
- explicitly created, used, released by apps
- client/server paradigm
- two types of transport service via socket API:
  - ❖ unreliable datagram
  - ❖ reliable, byte stream-oriented

### socket

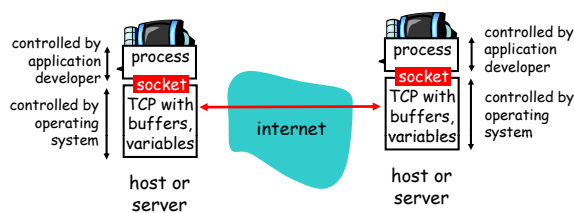
a *host-local, application-created, OS-controlled* interface (a "door") into which application process can **both send and receive** messages to/from another application process

2: Application Layer 92

## Socket-programming using TCP

**Socket:** a door between application process and end-end-transport protocol (UCP or TCP)

**TCP service:** reliable transfer of **bytes** from one process to another



2: Application Layer 93

## Socket programming *with TCP*

### Client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

### Client contacts server by:

- creating client-local TCP socket
- specifying IP address, port number of server process
- When **client creates socket**: client TCP establishes connection to server TCP

- When contacted by client, **server TCP creates new socket** for server process to communicate with client
  - ❖ allows server to talk with multiple clients
  - ❖ source port numbers used to distinguish clients (*more in Chap 3*)

### application viewpoint

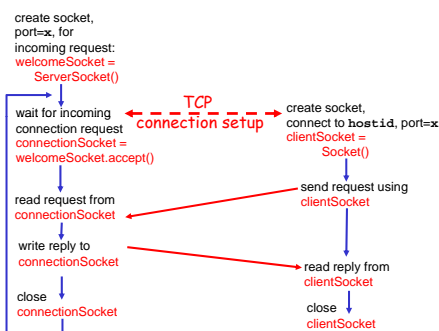
*TCP provides reliable, in-order transfer of bytes ("pipe") between client and server*

2: Application Layer 94

## Client/server socket interaction: TCP

Server (running on *hostid*)

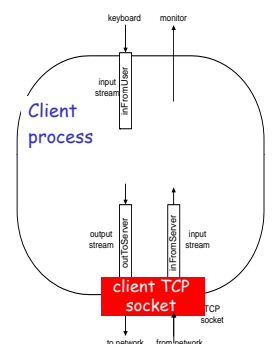
Client



2: Application Layer 95

## Stream jargon

- A **stream** is a sequence of characters that flow into or out of a process.
- An **input stream** is attached to some input source for the process, e.g., keyboard or socket.
- An **output stream** is attached to an output source, e.g., monitor or socket.



2: Application Layer 96

## Socket programming with TCP

### Example client-server app:

- 1) client reads line from standard input (`inFromUser` stream) , sends to server via socket (`outToServer` stream)
- 2) server reads line from socket
- 3) server converts line to uppercase, sends back to client
- 4) client reads, prints modified line from socket (`inFromServer` stream)

2: Application Layer 97

## Example: Java client (TCP)

```
import java.io.*;
import java.net.*;
class TCPClient {

    public static void main(String argv[]) throws Exception
    {
        String sentence;
        String modifiedSentence;

        Create input stream → BufferedReader inFromUser =
                                new BufferedReader(new InputStreamReader(System.in));

        Create client socket, connect to server → Socket clientSocket = new Socket("hostname", 6789);

        Create output stream attached to socket → DataOutputStream outToServer =
                                                    new DataOutputStream(clientSocket.getOutputStream());
```

2: Application Layer 98

## Example: Java client (TCP), cont.

```
        Create input stream attached to socket → BufferedReader inFromServer =
                                                new BufferedReader(new
                                                    InputStreamReader(clientSocket.getInputStream()));

        sentence = inFromUser.readLine();

        Send line to server → outToServer.writeBytes(sentence + '\n');

        Read line from server → modifiedSentence = inFromServer.readLine();
                               System.out.println("FROM SERVER: " + modifiedSentence);

        clientSocket.close();
    }
}
```

2: Application Layer 99

## Example: Java server (TCP)

```
import java.io.*;
import java.net.*;

class TCPServer {

    public static void main(String argv[]) throws Exception
    {
        String clientSentence;
        String capitalizedSentence;

        Create welcoming socket at port 6789 → ServerSocket welcomeSocket = new ServerSocket(6789);

        while(true) {
            Wait, on welcoming socket for contact by client → Socket connectionSocket = welcomeSocket.accept();

            Create input stream, attached to socket → BufferedReader inFromClient =
                                                        new BufferedReader(new
                                                            InputStreamReader(connectionSocket.getInputStream()));
```

2: Application Layer 100

## Example: Java server (TCP), cont

```
        Create output stream, attached to socket → DataOutputStream outToClient =
                                                    new DataOutputStream(connectionSocket.getOutputStream());

        Read in line from socket → clientSentence = inFromClient.readLine();

        capitalizedSentence = clientSentence.toUpperCase() + '\n';

        Write out line to socket → outToClient.writeBytes(capitalizedSentence);
    }
}

End of while loop, loop back and wait for another client connection
```

2: Application Layer 101

## Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
  - ✦ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P applications
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP

2: Application Layer 102

## Socket programming *with UDP*

UDP: no "connection" between client and server

- no handshaking
- sender explicitly attaches IP address and port of destination to each packet
- server must extract IP address, port of sender from received packet

UDP: transmitted data may be received out of order, or lost

application viewpoint

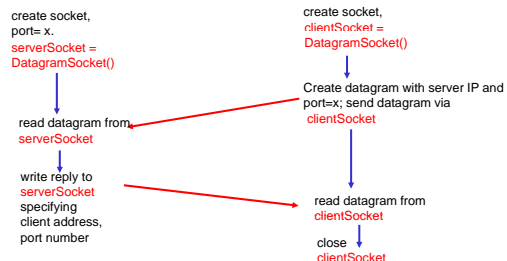
UDP provides *unreliable* transfer of groups of bytes ("datagrams") between client and server

2: Application Layer 103

## Client/server socket interaction: UDP

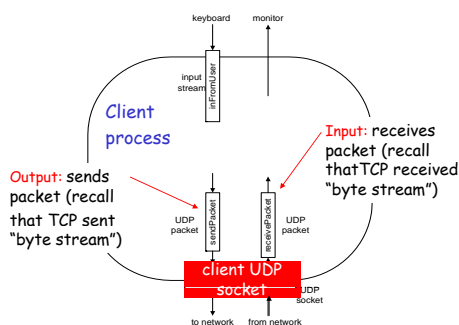
Server (running on hostia)

Client



2: Application Layer 104

## Example: Java client (UDP)



2: Application Layer 105

## Example: Java client (UDP)

```

import java.io.*;
import java.net.*;

class UDPClient {
    public static void main(String args[]) throws Exception {
        // Create input stream
        BufferedReader inFromUser =
            new BufferedReader(new InputStreamReader(System.in));

        // Create client socket
        DatagramSocket clientSocket = new DatagramSocket();

        // Translate hostname to IP address using DNS
        InetAddress IPAddress = InetAddress.getByName("hostname");

        byte[] sendData = new byte[1024];
        byte[] receiveData = new byte[1024];

        String sentence = inFromUser.readLine();
        sendData = sentence.getBytes();
  
```

2: Application Layer 106

## Example: Java client (UDP), cont.

```

        // Create datagram with data-to-send, length, IP addr, port
        DatagramPacket sendPacket =
            new DatagramPacket(sendData, sendData.length, IPAddress, 9876);

        // Send datagram to server
        clientSocket.send(sendPacket);

        DatagramPacket receivePacket =
            new DatagramPacket(receiveData, receiveData.length);

        // Read datagram from server
        clientSocket.receive(receivePacket);

        String modifiedSentence =
            new String(receivePacket.getData());

        System.out.println("FROM SERVER:" + modifiedSentence);
        clientSocket.close();
    }
  
```

2: Application Layer 107

## Example: Java server (UDP)

```

import java.io.*;
import java.net.*;

class UDPServer {
    public static void main(String args[]) throws Exception {
        // Create datagram socket at port 9876
        DatagramSocket serverSocket = new DatagramSocket(9876);

        byte[] receiveData = new byte[1024];
        byte[] sendData = new byte[1024];

        while(true) {
            // Create space for received datagram
            DatagramPacket receivePacket =
                new DatagramPacket(receiveData, receiveData.length);

            // Receive datagram
            serverSocket.receive(receivePacket);
  
```

2: Application Layer 108

## Example: Java server (UDP), cont

```
String sentence = new String(receivePacket.getData());

// Get IP addr, port #, of sender
InetAddress IPAddr = receivePacket.getAddress();
int port = receivePacket.getPort();

String capitalizedSentence = sentence.toUpperCase();

sendData = capitalizedSentence.getBytes();

// Create datagram to send to client
DatagramPacket sendPacket =
    new DatagramPacket(sendData, sendData.length, IPAddr,
        port);

// Write out datagram to socket
serverSocket.send(sendPacket);
}
```

End of while loop,  
loop back and wait for  
another datagram

2: Application Layer 109

## Chapter 2: Summary

our study of network apps now complete!

- application architectures
    - ❖ client-server
    - ❖ P2P
    - ❖ hybrid
  - application service requirements:
    - ❖ reliability, bandwidth, delay
  - Internet transport service model
    - ❖ connection-oriented, reliable: TCP
    - ❖ unreliable, datagrams: UDP
  - specific protocols:
    - ❖ HTTP
    - ❖ FTP
    - ❖ SMTP, POP, IMAP
    - ❖ DNS
    - ❖ P2P: BitTorrent, Skype
  - socket programming
- 2: Application Layer 110

## Chapter 2: Summary

Most importantly: learned about *protocols*

- typical request/reply message exchange:
    - ❖ client requests info or service
    - ❖ server responds with data, status code
  - message formats:
    - ❖ headers: fields giving info about data
    - ❖ data: info being communicated
- Important themes:*
- control vs. data msgs
    - ❖ in-band, out-of-band
  - centralized vs. decentralized
  - stateless vs. stateful
  - reliable vs. unreliable msg transfer
  - "complexity at network edge"
- 2: Application Layer 111